

Applied Class #1 - R Refresher

Introduction

This applied class will give you an opportunity to practice your R skills. For background material on R programming see my [crash course \[solutions\]](#) along with my introduction to [dplyr \[solutions\]](#), and my introduction to [ggplot \[solutions\]](#).

Part 1 - Discrimination in the Labor Market

In this question you'll partially replicate a well-known paper on racial bias in the labor market: "[Are Emily and Greg More Employable Than Lakisha and Jamal? A Field Experiment on Labor Market Discrimination](#)" by Marianne Bertrand and Sendhil Mullainathan. The paper, which I'll refer to as BM for short, appears in Volume 94, Issue #4 of the *American Economic Review*. You will need to consult this paper to complete this problem.

For convenience, I've posted a copy of the dataset from this paper on my website at https://ditraglia.com/data/lakisha_aer.csv. Each row of the dataset corresponds to a single fictitious job applicant. After loading the tidyverse library, you can read the data into a tibble called `bm` using the `read_csv()` function as follows:

```
library(tidyverse)
bm <- read_csv('https://ditraglia.com/data/lakisha_aer.csv')
```

You may want to consult the introduction and conclusion of BM, along with parts A-D of Section II of their paper as you work through the following exercises.

- Carry out the following steps:
 - Display the tibble `bm`. How many rows and columns does it have?
 - Display only the columns `sex`, `race` and `firstname` of `bm`. What information do these columns contain? How are `sex` and `race` encoded?
 - Add two new columns to `bm`: `female` should take the value `TRUE` if `sex` is female, and `black` should take value `TRUE` if `race` is black.
- Randomized controlled trials are all about *balance*: when the treatment is randomly assigned, the characteristics of the treatment and control groups will be the same on average. To answer the following parts you'll need a few additional pieces of information. First, the variable `computerskills` takes on the value `1` if a given resume says that the applicant has computer skills. Second, the variables `education` and `yearsexp` indicate level of education and years experience, while `ofjobs` indicates the number of previous jobs listed on the resume.
 - Is sex balanced across race? Use `dplyr` to answer this question. Hint: what happens if you apply the function `sum` to a vector of `TRUE` and `FALSE` values?
 - Are computer skills balanced across race? Hint: the summary statistic you'll want to use is the *proportion* of individuals in each group with computer skills. If you have a vector of ones and zeros, there is a very easy way to compute this.
 - Are `education` and `ofjobs` balanced across race?
 - Compute the mean and standard deviation of `yearsexp` by race. Comment on your findings.

- Why do we care if `sex`, `education`, `ofjobs`, `computerskills`, and `yearsexp` are balanced across race?
- Is `computerskills` balanced across `sex`? What about `education`? What's going on here? Is it a problem? Hint: re-read section II C of the paper.
- The outcome of interest in `bm` is `call` which takes on the value `1` if the corresponding resume elicits an email or telephone callback for an interview. Check your answers to the following against Table 1 of the paper:
 - Calculate the average callback rate for all resumes in `bm`.
 - Calculate the average callback rates separately for resumes with "white-sounding" and "black-sounding" names. What do your results suggest?
 - Repeat part 2, but calculate the average rates for each combination of race and sex. What do your results suggest?
- Read the help files for the `dplyr` function `pull()` and the base R function `t.test()`. Then test the null hypothesis that there is no difference in callback rates across black and white-sounding names against the two-sided alternative. Comment on your results.

Part 2 - Making Change

Do you remember physical money? In olden times I used to collect small change in a jar and periodically sort through it so I could do my laundry. I always seemed to have too many of the coins I *didn't* need, and to few of the ones I *did* need! This problem is about making change; it's also about making *functions*. To answer it you'll need two pieces of outside information. First: US coinage consists of pennies (\$0.01), nickels (\$0.05), dimes (\$0.10), quarters (\$0.25), and dollar coins (\$1.00). Second: UK coinage consists of pence, 2p coins, 5p coins, 10p coins, 20p coins, 50p coins, 1 pound coins, and 2 pound coins.

You'll also need to use the R operators `%/%` and `%%`. These represent *integer division* and *modulo*. For example `16` divided by `3` equals `5` with a remainder of `1` so `16 %/% 3` returns `5` while `16 %% 3` returns `1`. Before beginning this exercise, experiment with `%/%` and `%%` to make sure you understand how they work. If necessary, search the internet to find out more about them or read the R help file "Arithmetic Operators."

For technical reasons that I won't delve into here, the simplest reliable way of representing monetary values in a computer is by using *whole numbers* rather than decimals. Rather than representing fourteen pounds and thirty-two pence as `14.32`, for example, we'd store this value as `1432`. In short, we store the number of *pence* (or cents) as an *integer*. Please follow this convention throughout.

- The following code block outlines a simple algorithm for making change. The customer is owed a certain number of cents, an integer between `1` and `99`, and your task is to calculate how many quarters, dimes, nickels, and pennies to remove from the cash register. The rule is to always give as many *high denomination* coins as you can before moving on to lower denomination coins. If the customer is owed \$0.75, for example, you remove three quarters rather than seven dimes and a nickel. Fill in the gaps in the code, run it, and display the resulting vector `change`. Check the result by hand to make sure your logic is correct.

```
cents <- 73
quarters <- cents %/% 25
cents <- cents %% 25
```

```
dimes <- _____ # Delete the underscores & fill in the gap!
cents <- _____ # ditto!

nickels <- _____ # ditto!
cents <- _____ # ditto!

change <- c('quarters' = quarters,
            'dimes' = dimes,
            'nickels' = nickels,
            'pennies' = cents)
```

2. You've written some handy code! But now suppose you wanted to reuse it to make change for a different amount, say \$0.37 or \$0.19. Copying-and-pasting your existing code is tedious and error prone. Instead of doing this, write a *function* called `make_change_US()`. It should take a single input argument `cents`, a integer between 1 and 99, and return a vector called `change` listing the number of quarters, dimes, nickels, and pennies formatted as above. Run your function to make change for \$0.37 and \$0.19 and check the results against your hand calculations.

3. The function `make_change_US()` has some limitations. Suppose we wanted to make change for *UK* currency rather than US currency. This would require us to write a completely new function despite the underlying logic of the problem being identical. There's a better way forward. Write a new function called `make_change()` that it takes *two* input arguments. The first argument is `cents`. As before, this is a integer between 1 and 99. The second argument is new: `currency` is a named vector that stores the denominations between 1 and 99 cents along with their labels. For example we would set `currency` equal to `c('quarter' = 25, 'dime' = 10, 'nickel' = 5, 'penny' = 1)` for US currency and `c('50p' = 50, '20p' = 20, '10p' = 10, '5p' = 5, '2p' = 2, '1p' = 1)` for UK currency. As above, `make_change()` should return a named vector called `change` indicating how many of each coin to give as change, but now this vector should take its names from `currency`. The logic is the same as above, but the implementation is a bit trickier. While there are various ways to solve this, the simplest is probably to use a `for` loop. Test `make_change()` by making change for 78 cents/pence in US/UK currency. Compare your results to `make_change_US(78)` to double-check.