

# Lecture 003

## Resampling

---

Edward Rubin

Admin



# Admin

## Class today

### Review

- Regression and loss
- Classification
- KNN
- The bias-variance tradeoff

### Resampling methods

- Cross validation 
- The bootstrap 

# Admin

## Upcoming

### **Readings**

Today: *ISL* Ch. 5

**Problem set** Coming very soon...

Review

# Review

## Regression and loss

For **regression settings**, the loss is our **prediction**'s distance from **truth**, i.e.,

$$\text{error}_i = y_i - \hat{y}_i \quad \text{loss}_i = |y_i - \hat{y}_i| = |\text{error}_i|$$

Depending upon our ultimate goal, we choose **loss/objective functions**.

$$\text{L1 loss} = \sum_i |y_i - \hat{y}_i|$$

$$\text{MAE} = \frac{1}{n} \sum_i |y_i - \hat{y}_i|$$

$$\text{L2 loss} = \sum_i (y_i - \hat{y}_i)^2$$

$$\text{MSE} = \frac{1}{n} \sum_i (y_i - \hat{y}_i)^2$$

Whatever we're using, we care about **test performance** (e.g., test MSE), rather than training performance.

# Review

## Classification

For **classification problems**, we often use the **test error rate**.

$$\frac{1}{n} \sum_{i=1}^n \mathbb{I}(y_i \neq \hat{y}_i)$$

The **Bayes classifier**

1. predicts class  $j$  when  $\Pr(y_0 = j | \mathbf{X} = \mathbf{x}_0)$  exceeds all other classes.
2. produces the **Bayes decision boundary**—the decision boundary with the lowest test error rate.
3. is unknown: we must predict  $\Pr(y_0 = j | \mathbf{X} = \mathbf{x}_0)$ .

# Review

## KNN

**K-nearest neighbors** (KNN) is a non-parametric method for estimating

$$\Pr(y_0 = j | \mathbf{X} = \mathbf{x}_0)$$

that makes a prediction using the most-common class among an observation's "nearest" K neighbors.

- **Low values of K** (e.g., 1) are extremely flexible but tend to overfit (increase variance).
- **Large values of K** (e.g., N) are very inflexible—essentially making the same prediction for each observation.

The *optimal* value of K will trade off between overfitting and accuracy.



# Review

## The bias-variance tradeoff

Finding the optimal level of flexibility highlights the **bias-variance tradeoff**.

**Bias** The error that comes from inaccurately estimating  $f$ .

- More flexible models are better equipped to recover complex relationships ( $f$ ), reducing bias. (Real life is seldom linear.)
- Simpler (less flexible) models typically increase bias.

**Variance** The amount  $\hat{f}$  would change with a different **training sample**

- If new **training sets** drastically change  $\hat{f}$ , then we have a lot of uncertainty about  $f$  (and, in general,  $\hat{f} \neq f$ ).
- More flexible models generally add variance to  $f$ .

# Review

## The bias-variance tradeoff

The expected value<sup>†</sup> of the **test MSE** can be written


$$E \left[ \left( \mathbf{y}_0 - \hat{f}(\mathbf{X}_0) \right)^2 \right] = \underbrace{\text{Var} \left( \hat{f}(\mathbf{X}_0) \right)}_{\text{Variance}} + \underbrace{\left[ \text{Bias} \left( \hat{f}(\mathbf{X}_0) \right) \right]^2}_{\text{Bias}} + \underbrace{\text{Var}(\varepsilon)}_{\text{Irr. error}}$$

**The tradeoff** in terms of model flexibility

- Increasing flexibility *from total inflexibility* generally **reduces bias more** than it increases variance (reducing test MSE).
- At some point, the marginal benefits of flexibility **equal** marginal costs.
- Past this point (optimal flexibility), we **increase variance more** than we reduce bias (increasing test MSE).

**U-shaped test MSE** with respect to model flexibility (KNN here).

Increases in variance eventually overcome reductions in (squared) bias.



# Resampling methods

# Resampling methods

## Intro

**Resampling methods** help understand uncertainty in statistical modeling.

# Resampling methods

## Intro

**Resampling methods** help understand uncertainty in statistical modeling.

- *Ex. Linear regression:* How precise is your  $\hat{\beta}_1$ ?
- *Ex. With KNN:* Which K minimizes (out-of-sample) test MSE?

# Resampling methods

## Intro

**Resampling methods** help understand uncertainty in statistical modeling.

- *Ex. Linear regression:* How precise is your  $\hat{\beta}_1$ ?
- *Ex. With KNN:* Which K minimizes (out-of-sample) test MSE?

The process behind the magic of resampling methods:

1. **Repeatedly draw samples** from the **training data**.
2. **Fit your model(s)** on each random sample.
3. **Compare** model performance (or estimates) **across samples**.
4. Infer the **variability/uncertainty in your model** from (3).

# Resampling methods

## Intro

**Resampling methods** help understand uncertainty in statistical modeling.

- *Ex. Linear regression:* How precise is your  $\hat{\beta}_1$ ?
- *Ex. With KNN:* Which K minimizes (out-of-sample) test MSE?

The process behind the magic of resampling methods:

1. **Repeatedly draw samples** from the **training data**.
2. **Fit your model(s)** on each random sample.
3. **Compare** model performance (or estimates) **across samples**.
4. Infer the **variability/uncertainty in your model** from (3).

*Warning<sub>1</sub>* Resampling methods can be computationally intensive.

*Warning<sub>2</sub>* Certain methods don't work in certain settings.



# Resampling methods

## Today

Let's distinguish between two important **modeling tasks**:

- **Model selection** Choosing and tuning a model
- **Model assessment** Evaluating a model's accuracy

# Resampling methods

## Today

Let's distinguish between two important **modeling tasks**:

- **Model selection** Choosing and tuning a model
- **Model assessment** Evaluating a model's accuracy

We're going to focus on two common **resampling methods**:

1. **Cross validation** used to estimate test error, evaluating performance or selecting a model's flexibility
2. **Bootstrap** used to assess accuracy—parameter estimates or methods

# Resampling methods

## Hold out

*Recall:* We want to find the model that **minimizes out-of-sample test error**.

If we have a large test dataset, we can use it (once).

Q<sub>1</sub> What if we don't have a test set?

Q<sub>2</sub> What if we need to select and train a model?

Q<sub>3</sub> How can we avoid overfitting our training<sup>†</sup> data during model selection?

<sup>†</sup> Also relevant for *testing* data.

# Resampling methods

## Hold out

*Recall:* We want to find the model that **minimizes out-of-sample test error**.

If we have a large test dataset, we can use it (once).

Q<sub>1</sub> What if we don't have a test set?

Q<sub>2</sub> What if we need to select and train a model?

Q<sub>3</sub> How can we avoid overfitting our training<sup>†</sup> data during model selection?

A<sub>1,2,3</sub> **Hold-out methods** (e.g., cross validation) use training data to estimate test performance—**holding out** a mini "test" sample of the training data that we use to estimate the test error.

<sup>†</sup> Also relevant for *testing* data.

# Hold-out methods

## Option 1: The *validation set* approach

To estimate the **test error**, we can *hold out* a subset of our **training data** and then **validate** (evaluate) our model on this held out **validation set**.


- The **validation error rate** estimates the **test error rate**
- The model only "sees" the non-validation subset of the **training data**.

# Hold-out methods

## Option 1: The *validation set* approach

To estimate the **test error**, we can *hold out* a subset of our **training data** and then **validate** (evaluate) our model on this held out **validation set**.

- The **validation error rate** estimates the **test error rate**
- The model only "sees" the non-validation subset of the **training data**.




**Initial training set**

# Hold-out methods

## Option 1: The *validation set* approach

To estimate the **test error**, we can *hold out* a subset of our **training data** and then **validate** (evaluate) our model on this held out **validation set**.

- The **validation error rate** estimates the **test error rate**
- The model only "sees" the non-validation subset of the **training data**.



**Validation (sub)set**


**Training set:** Model training

# Hold-out methods

## Option 1: The *validation set* approach

To estimate the **test error**, we can *hold out* a subset of our **training data** and then **validate** (evaluate) our model on this held out **validation set**.

- The **validation error rate** estimates the **test error rate**
- The model only "sees" the non-validation subset of the **training data**.



**Validation (sub)set**

**Training set:** Model training



# Hold-out methods

## Option 1: The *validation set* approach

*Example* We could use the validation-set approach to help select the degree of a polynomial for a linear-regression model ([Kaggle](#)).

# Hold-out methods

## Option 1: The *validation set* approach

*Example* We could use the validation-set approach to help select the degree of a polynomial for a linear-regression model (Kaggle).

The goal of the validation set is to **estimate out-of-sample (test) error**.

Q So what?

# Hold-out methods

## Option 1: The *validation set* approach

*Example* We could use the validation-set approach to help select the degree of a polynomial for a linear-regression model (Kaggle).

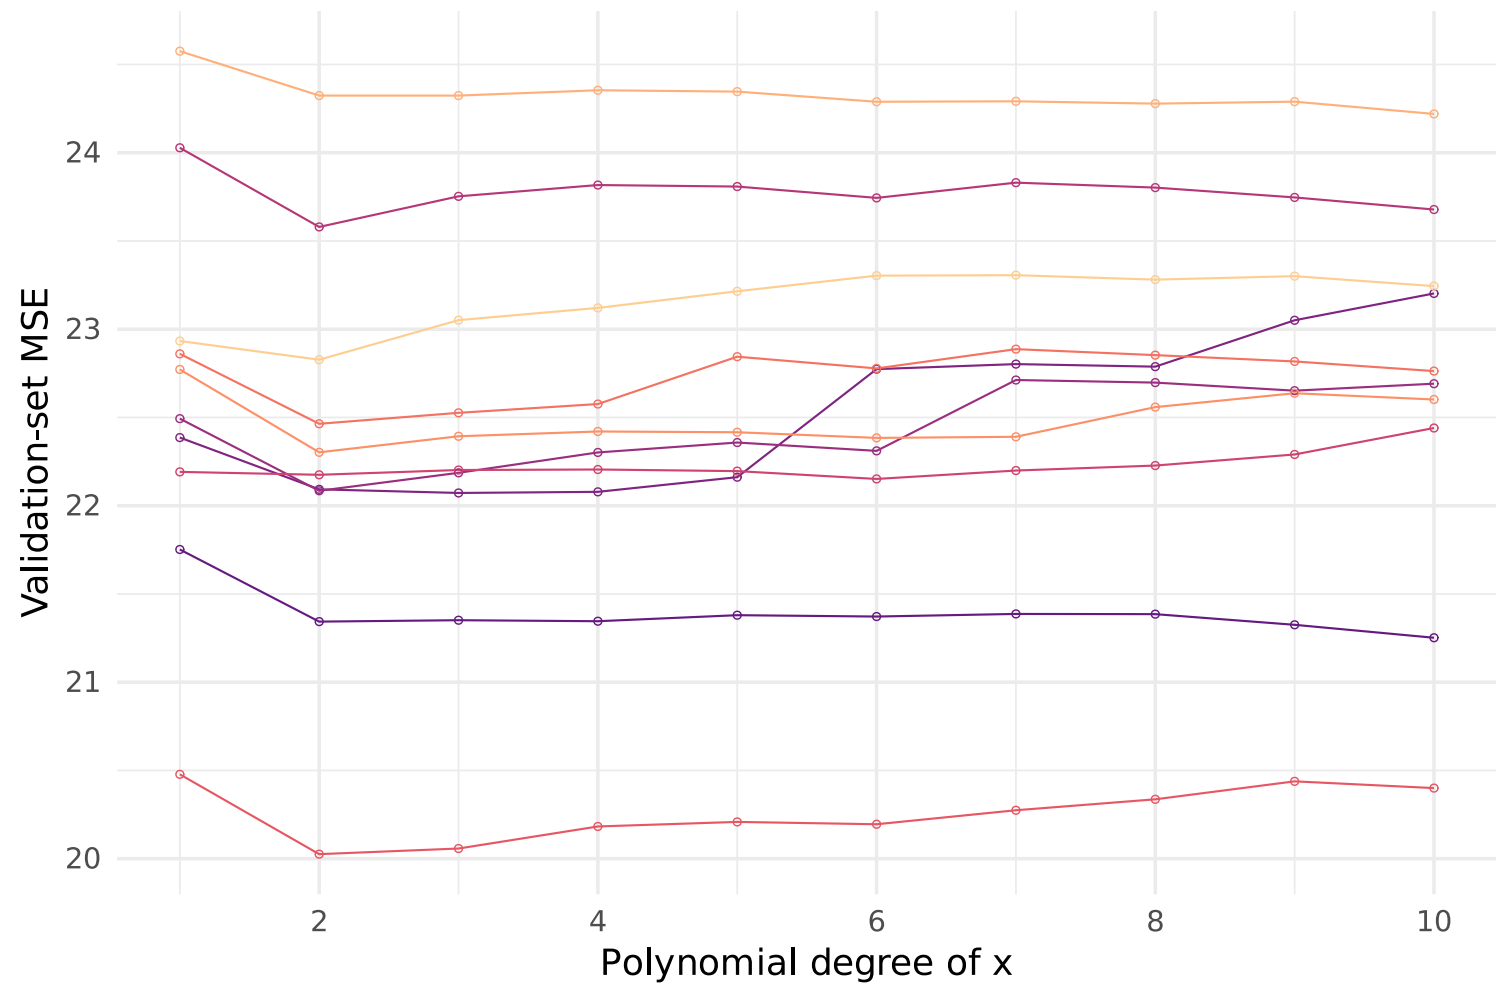
The goal of the validation set is to **estimate out-of-sample (test) error**.

Q So what?


- Estimates come with **uncertainty**—varying from sample to sample.
- Variability (standard errors) is larger with **smaller samples**.

**Problem** This estimated error is often based upon a fairly small sample (<30% of our training data). So its variance can be large.

**Validation MSE** for 10 different validation samples



**True test MSE** compared to validation-set estimates



# Hold-out methods

## Option 1: The *validation set* approach

Put differently: The validation-set approach has ( $\geq$ ) two major drawbacks:

1. **High variability** Which observations are included in the validation set can greatly affect the validation MSE.
2. **Inefficiency in training our model** We're essentially throwing away the validation data when training the model—"wasting" observations.

# Hold-out methods

## Option 1: The *validation set* approach

Put differently: The validation-set approach has ( $\geq$ ) two major drawbacks:

1. **High variability** Which observations are included in the validation set can greatly affect the validation MSE.
2. **Inefficiency in training our model** We're essentially throwing away the validation data when training the model—"wasting" observations.

(2)  $\Rightarrow$  validation MSE may overestimate test MSE.

Even if the validation-set approach provides an unbiased estimator for test error, it is likely a pretty noisy estimator.

# Hold-out methods

## Option 2: Leave-one-out cross validation

**Cross validation** solves the validation-set method's main problems.

- Use more (= all) of the data for training (lower variability; less bias).
- Still maintains separation between training and validation subsets.



# Hold-out methods

## Option 2: Leave-one-out cross validation

**Cross validation** solves the validation-set method's main problems.

- Use more (= all) of the data for training (lower variability; less bias).
- Still maintains separation between training and validation subsets.


**Leave-one-out cross validation** (LOOCV) is perhaps the cross-validation method most similar to the validation-set approach.

- Your validation set is exactly one observation.
- *New* You repeat the validation exercise for every observation.
- *New* Estimate MSE as the mean across all observations.

# Hold-out methods

## Option 2: Leave-one-out cross validation

Each observation takes a turn as the **validation set**, while the other  $n-1$  observations get to **train the model**.




Observation 1's turn for validation produces  $MSE_1$ .

# Hold-out methods

## Option 2: Leave-one-out cross validation

Each observation takes a turn as the **validation set**, while the other  $n-1$  observations get to **train the model**.




Observation 2's turn for validation produces  $MSE_2$ .

# Hold-out methods

## Option 2: Leave-one-out cross validation

Each observation takes a turn as the **validation set**, while the other  $n-1$  observations get to **train the model**.




Observation 3's turn for validation produces  $MSE_3$ .

# Hold-out methods

## Option 2: Leave-one-out cross validation

Each observation takes a turn as the **validation set**, while the other  $n-1$  observations get to **train the model**.




Observation 4's turn for validation produces  $MSE_4$ .

# Hold-out methods

## Option 2: Leave-one-out cross validation

Each observation takes a turn as the **validation set**, while the other  $n-1$  observations get to **train the model**.




Observation 5's turn for validation produces  $MSE_5$ .

# Hold-out methods

## Option 2: Leave-one-out cross validation

Each observation takes a turn as the **validation set**, while the other  $n-1$  observations get to **train the model**.



Observation  $n$ 's turn for validation produces  $MSE_n$ .

# Hold-out methods

## Option 2: Leave-one-out cross validation

Because **LOOCV uses  $n-1$  observations** to train the model,<sup>†</sup>  $MSE_i$  (validation MSE from observation  $i$ ) is approximately unbiased for test MSE.

**Problem**  $MSE_i$  is a terribly noisy estimator for test MSE (albeit  $\approx$ unbiased).

<sup>†</sup> And because often  $n-1 \approx n$ .



# Hold-out methods

## Option 2: Leave-one-out cross validation

Because **LOOCV uses n-1 observations** to train the model,<sup>†</sup>  $\text{MSE}_i$  (validation MSE from observation  $i$ ) is approximately unbiased for test MSE.

**Problem**  $\text{MSE}_i$  is a terribly noisy estimator for test MSE (albeit  $\approx$ unbiased).

**Solution** Take the mean!

$$\text{CV}_{(n)} = \frac{1}{n} \sum_{i=1}^n \text{MSE}_i$$

<sup>†</sup> And because often  $n-1 \approx n$ .

# Hold-out methods

## Option 2: Leave-one-out cross validation

Because **LOOCV uses n-1 observations** to train the model,<sup>†</sup>  $\text{MSE}_i$  (validation MSE from observation  $i$ ) is approximately unbiased for test MSE.

**Problem**  $\text{MSE}_i$  is a terribly noisy estimator for test MSE (albeit  $\approx$ unbiased).


**Solution** Take the mean!

$$\text{CV}_{(n)} = \frac{1}{n} \sum_{i=1}^n \text{MSE}_i$$

1. LOOCV **reduces bias** by using  $n-1$  (almost all) observations for training.
2. LOOCV **resolves variance**: it makes all possible comparisons (no dependence upon which validation-test split you make).

<sup>†</sup> And because often  $n-1 \approx n$ .

## True test MSE and LOOCV MSE compared to validation-set estimates



# Hold-out methods

## Option 3: k-fold cross validation

Leave-one-out cross validation is a special case of a broader strategy:

### **k-fold cross validation.**

1. **Divide** the training data into  $k$  equally sized groups (folds).
2. **Iterate** over the  $k$  folds, treating each as a validation set once (training the model on the other  $k - 1$  folds).
3. **Average** the folds' MSEs to estimate test MSE.

# Hold-out methods

## Option 3: k-fold cross validation

Leave-one-out cross validation is a special case of a broader strategy:  
**k-fold cross validation.**

1. **Divide** the training data into  $k$  equally sized groups (folds).
2. **Iterate** over the  $k$  folds, treating each as a validation set once (training the model on the other  $k - 1$  folds).
3. **Average** the folds' MSEs to estimate test MSE.

Benefits?

# Hold-out methods

## Option 3: k-fold cross validation

Leave-one-out cross validation is a special case of a broader strategy:  
**k-fold cross validation.**

1. **Divide** the training data into  $k$  equally sized groups (folds).
2. **Iterate** over the  $k$  folds, treating each as a validation set once (training the model on the other  $k - 1$  folds).
3. **Average** the folds' MSEs to estimate test MSE.

Benefits?

1. **Less computationally demanding** (fit model  $k = 5$  or  $10$  times; not  $n$ ).

# Hold-out methods

## Option 3: k-fold cross validation

Leave-one-out cross validation is a special case of a broader strategy:  
**k-fold cross validation.**

1. **Divide** the training data into  $k$  equally sized groups (folds).
2. **Iterate** over the  $k$  folds, treating each as a validation set once (training the model on the other  $k - 1$  folds).
3. **Average** the folds' MSEs to estimate test MSE.

Benefits?

1. **Less computationally demanding** (fit model  $k = 5$  or  $10$  times; not  $n$ ).
2. **Greater accuracy** (in general) due to bias-variance tradeoff!

# Hold-out methods

## Option 3: k-fold cross validation

Leave-one-out cross validation is a special case of a broader strategy:

### **k-fold cross validation.**

1. **Divide** the training data into  $k$  equally sized groups (folds).
2. **Iterate** over the  $k$  folds, treating each as a validation set once (training the model on the other  $k - 1$  folds).
3. **Average** the folds' MSEs to estimate test MSE.

Benefits?

1. **Less computationally demanding** (fit model  $k = 5$  or  $10$  times; not  $n$ ).
2. **Greater accuracy** (in general) due to bias-variance tradeoff!
  - Somewhat higher bias, relative to LOOCV:  $n - 1$  vs.  $(k - 1)/k$ .



# Hold-out methods

## Option 3: k-fold cross validation

Leave-one-out cross validation is a special case of a broader strategy:

### **k-fold cross validation.**

1. **Divide** the training data into  $k$  equally sized groups (folds).
2. **Iterate** over the  $k$  folds, treating each as a validation set once (training the model on the other  $k - 1$  folds).
3. **Average** the folds' MSEs to estimate test MSE.

Benefits?

1. **Less computationally demanding** (fit model  $k = 5$  or  $10$  times; not  $n$ ).
2. **Greater accuracy** (in general) due to bias-variance tradeoff!
  - Somewhat higher bias, relative to LOOCV:  $n - 1$  vs.  $(k - 1)/k$ .
  - Lower variance due to high-degree of correlation in LOOCV  $\text{MSE}_i$ .

# Hold-out methods

## Option 3: k-fold cross validation

Leave-one-out cross validation is a special case of a broader strategy:

### k-fold cross validation.

1. **Divide** the training data into  $k$  equally sized groups (folds).
2. **Iterate** over the  $k$  folds, treating each as a validation set once (training the model on the other  $k - 1$  folds).
3. **Average** the folds' MSEs to estimate test MSE.

Benefits?


1. **Less computationally demanding** (fit model  $k = 5$  or  $10$  times; not  $n$ ).
2. **Greater accuracy** (in general) due to bias-variance tradeoff!
  - Somewhat higher bias, relative to LOOCV:  $n - 1$  vs.  $(k - 1)/k$ .
  - Lower variance due to high-degree of correlation in LOOCV MSE. 🤔

# Hold-out methods

## Option 3: k-fold cross validation

With  $k$ -fold cross validation, we estimate test MSE as

$$CV_{(k)} = \frac{1}{k} \sum_{i=1}^k \text{MSE}_i$$




Our  $k = 5$  folds.

# Hold-out methods

## Option 3: k-fold cross validation

With  $k$ -fold cross validation, we estimate test MSE as

$$CV_{(k)} = \frac{1}{k} \sum_{i=1}^k \text{MSE}_i$$




Each fold takes a turn at **validation**. The other  $k - 1$  folds **train**.

# Hold-out methods

## Option 3: k-fold cross validation

With  $k$ -fold cross validation, we estimate test MSE as

$$CV_{(k)} = \frac{1}{k} \sum_{i=1}^k \text{MSE}_i$$




For  $k = 5$ , fold number 1 as the **validation set** produces  $\text{MSE}_{k=1}$ .

# Hold-out methods

## Option 3: k-fold cross validation

With  $k$ -fold cross validation, we estimate test MSE as

$$CV_{(k)} = \frac{1}{k} \sum_{i=1}^k \text{MSE}_i$$




For  $k = 5$ , fold number 2 as the **validation set** produces  $\text{MSE}_{k=2}$ .

# Hold-out methods

## Option 3: k-fold cross validation

With  $k$ -fold cross validation, we estimate test MSE as

$$CV_{(k)} = \frac{1}{k} \sum_{i=1}^k MSE_i$$




For  $k = 5$ , fold number 3 as the **validation set** produces  $MSE_{k=3}$ .

# Hold-out methods

## Option 3: k-fold cross validation

With  $k$ -fold cross validation, we estimate test MSE as

$$CV_{(k)} = \frac{1}{k} \sum_{i=1}^k \text{MSE}_i$$



For  $k = 5$ , fold number 4 as the **validation set** produces  $\text{MSE}_{k=4}$ .




# Hold-out methods

## Option 3: k-fold cross validation

With  $k$ -fold cross validation, we estimate test MSE as

$$CV_{(k)} = \frac{1}{k} \sum_{i=1}^k \text{MSE}_i$$



For  $k = 5$ , fold number 5 as the **validation set** produces  $\text{MSE}_{k=5}$ .


# Hold-out methods

## Option 3: k-fold cross validation

With  $k$ -fold cross validation, we estimate test MSE as

$$\text{CV}_{(k)} = \frac{1}{k} \sum_{i=1}^k \text{MSE}_i$$

**Test MSE** vs. estimates: LOOCV, 5-fold CV (20x), and validation set (10x)



*Note:* Each of these methods extends to classification settings, e.g., LOOCV

$$\text{CV}_{(n)} = \frac{1}{n} \sum_{i=1}^n \mathbb{I}(\mathbf{y}_i \neq \hat{\mathbf{y}}_i)$$

# Hold-out methods

## Caveat

So far, we've treated each observation as separate/independent from each other observation.

The methods that we've defined assume this **independence**.

# Hold-out methods

## Caveat

So far, we've treated each observation as separate/independent from each other observation.

The methods that we've defined assume this **independence**.

Make sure that you think about

- the **structure** of your data
- the **goal** of the prediction exercise

*E.g.,*

1. Are you trying to predict the behavior of **existing** or **new** customers?
2. Are you trying to predict **historical** or **future** recessions?

The bootstrap

# The bootstrap

## Intro

The **bootstrap** is a resampling method often used to quantify the uncertainty (variability) underlying an estimator or learning method.

### Hold-out methods

- randomly divide the sample into training and validation subsets
- train and validate ("test") model on each subset/division

### Bootstrapping

- randomly samples **with replacement** from the original sample
- estimates model on each of the *bootstrap samples*



# The bootstrap

## Intro

Estimating an estimate's standard error involves assumptions and theory.<sup>†</sup>

There are times this derivation is difficult or even impossible, *e.g.*,

$$\text{Var}\left(\frac{\hat{\beta}_1}{1 - \hat{\beta}_2}\right)$$

The bootstrap can help in these situations.

Rather than deriving an estimator's variance, we use bootstrapped samples to build a distribution and then learn about the estimator's variance.

<sup>†</sup> Recall the standard-error estimator for OLS.

# Intuition

*Idea:* Bootstrapping builds a distribution for the estimate using the variability embedded in the training sample.


# The bootstrap

## Graphically

$Z$

7	8	9
4	5	6
1	2	3

$$\hat{\beta} = 0.653$$




# The bootstrap

## Graphically

$Z$

7	8	9
4	5	6
1	2	3


$$\hat{\beta} = 0.653$$



$Z^{\star 1}$

7	9	3
9	3	8
3	9	9

$$\hat{\beta} = -0.96$$




# The bootstrap

## Graphically

$Z$

7	8	9
4	5	6
1	2	3


$$\hat{\beta} = 0.653$$



$Z^{*1}$

7	9	3
9	3	8
3	9	9


$$\hat{\beta} = -0.96$$



$Z^{*2}$


9	7	5
5	7	9
4	1	7

$$\hat{\beta} = 0.968$$



# The bootstrap


## Graphically



# The bootstrap

Running this bootstrap 10,000 times

```
plan(multisession, workers = 10)
# Set a seed
set.seed(123)
# Run the simulation 1e4 times
boot_df <- future_map_dfr(
  # Repeat sample size 100 for 1e4 times
  rep(n, 1e4),
  # Our function
  function(n) {
    # Estimates via bootstrap
    est <- lm(y ~ x, data = z[sample(1:n, n, replace = T), ])
    # Return a tibble
    data.frame(int = est$coefficients[1], coef = est$coefficients[2])
  },
  # Let furrr know we want to set a seed
  .options = furrr_options(seed = TRUE)
)
```





# The bootstrap

## Comparison: Standard-error estimates

The **bootstrapped standard error** of  $\hat{\alpha}$  is the standard deviation of the  $\hat{\alpha}^{*b}$

$$\text{SE}_B(\hat{\alpha}) = \sqrt{\frac{1}{B} \sum_{b=1}^B \left( \hat{\alpha}^{*b} - \frac{1}{B} \sum_{\ell=1}^B \hat{\alpha}^{*\ell} \right)^2}$$

This 10,000-sample bootstrap estimates  $\text{S.E.}(\hat{\beta}_1) \approx 0.77$ .

# The bootstrap


## Comparison: Standard-error estimates

The **bootstrapped standard error** of  $\hat{\alpha}$  is the standard deviation of the  $\hat{\alpha}^{*b}$

$$\text{SE}_B(\hat{\alpha}) = \sqrt{\frac{1}{B} \sum_{b=1}^B \left( \hat{\alpha}^{*b} - \frac{1}{B} \sum_{\ell=1}^B \hat{\alpha}^{*\ell} \right)^2}$$

This 10,000-sample bootstrap estimates  $\text{S.E.}(\hat{\beta}_1) \approx 0.77$ .

If we go the old-fashioned OLS route, we estimate 0.673.



# Resampling

## Review

### Previous resampling methods

- Split data into **subsets**:  $n_v$  validation and  $n_t$  training ( $n_v + n_t = n$ ).
- Repeat estimation on each subset.
- Estimate the true test error (to help tune flexibility).

### Bootstrap

- Randomly samples from training data **with replacement** to generate  $B$  "samples", each of size  $n$ .
- Repeat estimation on each subset.
- Estimate the variance estimate using variability across  $B$  samples.

# Sources

These notes draw upon

- [An Introduction to Statistical Learning \(ISL\)](#)  
James, Witten, Hastie, and Tibshirani
- [Python Data Science Handbook](#)  
Jake VanderPlas

# Table of contents

## Admin

- Today
- Upcoming

## Review

- Regression and loss
- Classification
- KNN
- The bias-variance tradeoff

## Examples

- Validation-set simulation
- LOOCV MSE
- k-fold CV

## Resampling

- Intro
- Hold-out methods
  - Validation sets
  - LOOCV
  - k-fold cross validation
- The bootstrap
  - Intro
  - Graphically
  - Example

## Other

- Sources/references